

What the knowledge base shows that the marketing site could not

The 29 July board already read Blink off its marketing site and got the architecture right. This pass read their published help centre instead, which carries the thing screenshots cannot: the exact labels, in order, for how a Blink operator declares a job done — and what the product does with that declaration afterwards.

SOURCE	ALSO READ
help.blink.global — 12 collections, 10 articles	blink.global platform page, driven live
MOBBIN	HANDS-ON
6 probes, all negative — Blink is not indexed	None. Nothing here was driven in their product

Read this first: most of this was already known

This addendum exists because a fresh search rediscovered findings the 29 July board already carried. That is worth recording rather than hiding, because it says something about where the research lives: the board is not indexed anywhere a repo grep or a Mobbin query can reach it, so it was invisible to the search that went looking.


FINDING	STATUS AGAINST THE BOARD
Two independent status columns	Already on the board — named there as "Status is two dimensions, not one", with the full chip list.
Send Failed must be its own state	Already on the board — "a boolean will not do", tied to our comms gap.
No readiness engine, no proposed action	Already on the board — "It configures... that is an admin console."
The unified agenda as the multi-day model	Already on the board — and it is the reference model for the keystone gap.

FINDING	STATUS AGAINST THE BOARD
The amber-tinted "now" row	Already acted on — flagged as a trap, found in our own phase spine, repointed to --progress .
★ The operator completion loop, in their labels	New — knowledge-base sourced. Not visible in any screenshot.
★ AI auto-fill does not mark what it filled	New — and it silently drops data.
The analytics inventory, itemised	New — turns the board's posture claim into a specific, checkable absence.
<i>What the 1 August pass found, against what the 29 July board already said. Only the bottom group is new material.</i>	

Source warning carried forward from the board, and now partly lifted: that page drew only on marketing material and published screenshots. This pass adds Blink's own help centre — author-stated, and far more specific — but it is still documentation, not the running product. Nothing here was driven. Where the docs describe a sequence, the sequence is theirs; whether the shipped app matches is unverified.

The operator completion loop — the material the board could not reach

Their driver app is where a Blink job is actually declared finished. The help centre gives the labels in order, and the ordering is the finding.



Completion is declared, then committed Take

Mark As Complete declares; Finish Trip commits. Clock In demands a swipe rather than a tap. Friction is spent only on state changes that are expensive to reverse, and nowhere else in the flow.

OUR SIDE	Resolves settle in one gesture. That is correct for most of the queue and wrong for the few that move money or tell a third party something.
VERDICT	Worth taking for the small set of irreversible settles. The optional comment is the valuable half — it makes a completed item carry <i>why</i> , which is provenance we would otherwise have to reconstruct.
DO NEXT	Identify which resolves are one-way, and give only those the second beat.

The exception is a peer of the completion, and does not end the job

Take

Report a Problem sits at the same level as Mark As Complete — not buried, not a fallback — and forks to End Trip or Continue Trip . Reporting a problem does not force the job to end. The same adjacency appears on their check-in screen, where Check In is stacked directly above Flag Attendee .

VERDICT	This is the sharpest thing in the addendum. A surface that offers only resolve-or-ignore trains hosts to mark things done that are not done — which corrupts the readiness signal our whole product rests on.
DO NEXT	Wherever we offer a resolve, check whether "it went wrong but is still running" is expressible. Today it mostly is not.

The confirmation crosses the role boundary

Note

The driver finishes; *organizers are notified*. The satisfied ask travels back to whoever asked it, rather than sitting in the completer's own view.

VERDICT	Relevant when helper and vendor confirm-backs land. A confirmation that only the responder can see is not a confirmation.
---------	---

Three kinds of not-done — extending the board's finding

The board established that Send Failed must be distinct from Not Sent and Pending . What the ownership mapping adds is why: each of the three points at a different party, and only one of them is an ask a host can act on.

STATE	MEANING	WHO OWNS IT	IS IT A HERO-ASK?
Not Sent	we have not asked yet	the host	Yes
Pending	we asked, no answer	the other party	No — it is a wait
Send Failed	we tried, the mechanism broke	the system	No — it is a defect
Denied	a real no	the other party	No
Approved	done	—	—

The Invitation Status axis, with the party each state actually points at. Owner column is this pass's reading, not a Blink label.

Two details worth keeping. **Denied** and **Send Failed** share the red token and are separated by glyph — a real no against a broken pipe. And where invitation status does not yet apply, they print an explicit em dash rather than leaving the cell blank: blank reads as unknown, the dash reads as not applicable at this stage. That is the same instinct as our field-level honesty copy, applied inside a table cell.

Why this matters more than it looks

Acts on propose-don't-ask

A readiness engine that treats all three not-dones as "incomplete" will keep proposing actions the host cannot take. Proposing "chase the vendor" when the real state is *we never wrote to them* is a different instruction from proposing it when they have had the message for six days — and proposing anything at all is wrong when the send itself failed.

Two guards on the ask itself

They refuse to silently re-ask

"If any attendees have already been invited, a notification will appear. Choose whether to send them another invitation or exclude them."

The system knows an ask was already made and interrupts rather than duplicating it. Ours should run the same check at the propose

Bulk confirms before it commits

Save & Continue → review summary → Send .

The check sits ahead of the act, not only behind it. Worth remembering when outbound comms ship — the moment we can actually send is the moment this stops being theoretical.

boundary, before the hero ask offers something the host already did.

Smaller, and free: their sessions report capacity as **Empty · Filling Up · Almost Full · Full** rather than a percentage. Four named thresholds beat a number nobody calibrates — the same argument as our status labels never being colour alone.

What their analytics actually count

The board concluded from the marketing site that Blink configures rather than guides. The help centre lets that be stated as a specific, checkable inventory rather than an impression.

GROUP	METRICS
Badges	Printed vs. Unprinted · Badge Print Rate · Purchased vs. Total · Assigned vs. Purchased · Target Badge Sold
Sessions	Ticket Utilization Rate · Ticketed Session Attendance Rate · fill state (Empty / Filling Up / Almost Full / Full)
Commerce	Total Revenue · Total Number of Orders · Average Order Value · Orders (Completed vs. Abandoned)
Event	Total Attendees with Badges · Total Event Staff Count
Work outstanding	Nothing. No count, no ratio, no rollup of asks resolved against asks raised.
Every metric named in Blink's Statistics guide, grouped. The last row is the absence, not a metric.	

They count inventory, not readiness. Completion aggregates for things sold and printed; never for work owed. Their platform can report 400 badges printed and 12 orders abandoned, and cannot tell an organiser whether the event is ready.

Two structural gaps underneath it

- **Journeys and convoys carry no documented status.** The lifecycle exists on the driver side; the organiser-side journey object has no state machine in the docs, and the convoy article covers no delay, disruption or resequencing. Coordination is capacity-

The same defect we have

Their convoy model is a set of independent vehicles with no cross-vehicle sequencing — which is precisely what our own vendor audit found on our side: isolated cockpits, nothing that orders one against another.

Convergence on a weakness is still information. It suggests cross-item sequencing is genuinely

based only — *"vehicles remain available until they reach full capacity."*

- **Table Views has no completion concept.** Filters, saved views, a filter-count badge, a 2,000-row limit — but status is just another filterable column. No grouped done-bucket, no counts.

hard rather than merely unbuilt, and that shipping it would be a real differentiator rather than catching up.

Where we are ahead — and it is the AI layer

Their auto-fill does not mark what it filled

We lead

Blink's AI flight extraction instructs the user to *"be sure to review the auto-filled details to confirm that all information is correct"* — but the interface does not distinguish an AI-populated field from a hand-entered one. It asks for a review without marking what to review.

WORSE	It silently collapses connections: it <i>"automatically fills in the total travel hours, including any transits or connections, without listing them individually. The system captures only the start and end destinations."</i> Inference presented as fact, with data loss, inside a travel record an operations team will act on.
OUR SIDE	Provenance at 31 sites, a <code>*_SOURCES</code> registry, and a doctrine that forbids exactly this.
VERDICT	The board established we are ahead on posture. This is the first place where an enterprise competitor's AI is not merely thinner than ours but <i>less honest</i> . Do not trade it away chasing their table density.

How the event cohort shows a satisfied ask

Mobbin-sourced, and new to this research: what the consumer set actually renders at the moment an ask is answered. The pattern is consistent enough to be a rule.

APP	WHAT A SATISFIED ASK LEAVES BEHIND
Apple Invites	Roster grouped GOING (7) · MAYBE (1) · NOT GOING (1) , and every response carries the person's own sentence — <i>"I'll need to check my calendar, get back to you!"</i>
Partiful	Hearts that break progressively across Going / Maybe / Can't Go; the count sits centre-stage in the host action bar; RSVPs accrue into a reliable activity feed

APP	WHAT A SATISFIED ASK LEAVES BEHIND
GroupMe	Attendee tabs Going (1) · Not Going (0) · Pending (1) — the un-answered tracked as hard as the answers
BFF	Answered rows read 2 <i>going</i> ; unanswered rows keep a standing RSVP button. The CTA is the absence of an answer
Blinkist	Every toast is backed by a durable marker — a green <i>Finished</i> chip on the object, and a filter that counts it

Read off Mobbin flows for each app, 1 August 2026.

A satisfied ask does not produce a notification — it mutates the roster. Only Blinkist, which is not an event app, leans on toasts at all, and even there every toast has a permanent counterpart on the object.

*This closes a loop the board opened. Its first mechanism — "state is promoted to a named surface" — was read off Wise and Luma and marked "we do it, one level down." Blink's status columns and the cohort's rosters are the same finding from two more directions: **the leaders carry state as an identity; we carry it as a parameter.** Three independent sources now. The row that scores it, Feedback & affordance, sits at 8/10.*

Sources. *help.blink.global — 12 collections enumerated, 10 articles read in full (Registration, Transportation, Product Updates, Organizer App, plus the Send Invitations, Add Attendees, Table Views, Add a New Journey, Blink Drive, Admission Settings, AI Flight Extraction, Statistics and Convoys articles). blink.global platform page driven live in-browser. Mobbin flows for Apple Invites, Partiful, GroupMe, BFF, Luma and Blinkist, read 1 August 2026.*

Not claimed. *No hands-on with any Blink product. No account, no trial, no flow taxonomy — Blink is not on Mobbin and six probes confirmed it. The completion loop above is Blink's own documented sequence; whether the shipped app matches its documentation is unverified. Where the help centre was silent — Application Status values, journey lifecycle — that silence is reported as silence rather than filled in.*

Companion to the Mobbin Competitive Read gap board, 29–30 July 2026, which remains the primary document. Rendered in Studio Matte on the product's own carbon ramp, single-theme by decision.