

Multi-option on mobile — and the one rule every leader follows

Across every app read for this, not one wraps a pill row onto a second line. Not one. If pills are laying out badly, the fix is almost never better wrapping — it is choosing which of five well-established containers the options belong in, and that choice is decided by how many there are and how many axes they span.

READ

Careem · Binance · Hyundai Card · komoot · Jobber · Nike · Zip · Crypto.com · Yubo · Formula 1 · foodpanda · Meta AI · Booking.com · Wise · Journal · Azar · Linktree · KakaoTalk

PLUS

blink.global mobile frames


DEMOS BELOW

Live and scrollable — drag them

The universal rule: nobody wraps


Eighteen apps. Zero wrapped pill rows. Every one either scrolls horizontally, escalates to a sheet, or converts to a grid. **Wrapping is the thing that looks like the obvious answer and is used by none of them.**

✗ WHAT BREAKS — WRAPPED PILLS



Ragged right edge, unpredictable height, the control reflows when the label set changes, and a

✓ THE SAME OPTIONS, SCROLLED



One line, fixed height, and the mask at the right edge clips mid-pill — the clip is the affordance that says there is more. Drag it.

two-line filter bar eats the fold. This is the failure mode.

The five containers, ranked by option count

OPTIONS	AXES	CONTAINER	BEST IN SET
2–4	1	Segmented control / underline tabs — equal width, all visible	Nike, Zip
5–8	1	Scrolling pill rail with a clipped right edge	Careem, Formula 1, KakaoTalk
5–8	1, but must all be reachable	Scrolling rail + a  expander at the end	Hyundai Card
any	2+	Stacked labelled rails — one axis per row	Binance
9+	1	Bottom sheet, radio rows, grouped with dividers	komoot
9+ visual	1	2-column tile grid, uppercase labels	blink.global, Crypto.com, Nike
<i>Choose the container from the count and the number of axes, not from taste.</i>			

Binance — the best answer for multiple axes

Three dimensions of choice, stacked as three labelled rails. No modal, no wrapping, and every axis shows its current value at a glance.

What needs you

All

Decisions

Vendors

Money

T

When

Any time

This week

Before the day

C

Standing

All

Open

Waiting on someone

Settle

WHY IT WINS A single rail forces unrelated options into one line, so the user must read every pill to find the axis they care about. Splitting by axis means **each row is scanned once and skipped** — and the muted label costs 11px of height to remove that work entirely.

COST Vertical space. Three axes is roughly 96px. Worth it above a long list; wasteful above a short one.

Three more mechanisms worth taking

Anchor the left

Careem · Jobber

A fixed icon-pill that does **not** scroll away, holding the overflow control or the active-filter count. Everything else scrolls past it.

⚙ 2

Needs a decision

Vendors

Scroll the rail — the anchor stays. Always a route back to the full set, and it carries the count of what is currently on.

The escape hatch

★ Hyundai Card

Scrolling rail *plus* a trailing expander. Solves the real complaint about rails — that options past the edge are effectively invisible — without giving up the single line.

All

Vendors

Money

Guests

The count in the expander is the honest part — it says how much you cannot see, rather than leaving it to be discovered by dragging.

Grouped sheet for long lists

★ komoot

Past roughly nine options a rail stops working at any width. komoot's sheet groups by kind with a divider and puts an *All ...* row at the head of each group, so a broad pick never requires scrolling to find.

Everything



PEOPLE

All guest items



RSVPs

2 still out



Dietary needs



MONEY

All money items



Deposits due



Who owes what



DETAIL

The radio sits **right**, the label left — every sheet in the set does this, and it means the thumb reaches the control without covering the text it is choosing.

SELECTED

Row-tinted *and* radio filled. Never colour alone.

What blink.global does — and it is the sharpest lesson

Their mobile app has **no pill rail at all**. Eight browse categories become a 2-column tile grid with uppercase labels; the agenda is four rows; the tab bar is five icons. There is nothing to lay out badly because there is almost nothing to lay out.

Before choosing a container, question the count. A rail of nine filters is often a symptom that one list is doing the work of three surfaces — which is exactly what our own density re-audit found. The best pill layout is frequently four fewer pills.


Caveat: Blink's option count is low because it is an attendee app with no obligations, not because they solved a hard problem. Their restraint is worth copying; their circumstances are not ours.

Nike and blink.global, translated to our data

Both use a 2-column tile grid, and they differ on exactly one decision — where the text sits. That single difference decides whether either one can carry our labels.


NIKE'S SHAPE · TEXT BELOW THE IMAGE

All Needs you Waiting Settled




Where everyone stays

Lake house · 3 nights
Committed · \$1,394




Getting around

4 rides · 2 unassigned
Sat pickup 11:00



Food & shopping

Bayside Crab · deposit paid
Count due Friday




Guests & RSVPs

42 of 44 replied
2 still out


Our labels survive intact. Two muted lines beneath carry the state without a badge — the tile reads as a place with a status, which is what a Sections door actually is.

BLINK'S SHAPE · TEXT OVERLAID, UPPERCASE


BROWSE ALL




WHERE EVERYO...




GETTING AROUND




FOOD & SHOPPI...



GUESTS & RSVPS



RUN OF SHOW



MONEY

The top-left tile shows the failure: overlaid uppercase truncates our longest label. Blink gets away with it because theirs are single words — CONFERENCE, MUSIC, EXPO.

The one decision that separates them

Nike's is the one that fits us

Blink overlays a short uppercase word on the artwork with a scrim. Nike puts the text *underneath*, at normal case, over two lines.

WHY IT MATTERS

Our surface names are planner-language sentences — "Where everyone stays", "Getting around", "Food & shopping". Blink's container demands single words, and the only way to fit them is to invent jargon — STAYS , GROUND — which UX_06 forbids and which would undo the exact naming work already shipped.

SECOND REASON

Text-below has no contrast problem. Text-on-photo needs a scrim on every image and still risks failing WCAG against a bright one. We would be adding a legibility risk to buy a shorter label we do not want.

WHAT BLINK STILL WINS

Equal-size tiles. A grid is ragged-proof by construction — the layout cannot break when a label changes length, which is precisely the complaint that started this study.

The honest problem with both — for us

Read before adopting

Nike's and Blink's grids are **browse** surfaces. Nothing in either is raising a count, and neither has to tell the user which tile needs them today.

OUR SECTIONS DIRECTORY DOES	It is a door to every surface <i>and</i> it carries what each surface is raising. Put a numeric badge on a tile — as the amber 2 above shows — and the grid stops being calm; leave it off and the directory stops being useful.
THE RESOLUTION	Use Nike's two muted lines as the count. <i>"4 rides · 2 unassigned"</i> says what a badge would, in the register we already use, without spending a colour. That is the translation — not the tile shape, the decision to put state in prose beneath rather than a number on top.
WHERE IT BELONGS	The Sections directory, which is browse-shaped by definition — <i>"a door to EVERY surface"</i> . Not the home fold, where the ranked ask lives.

Both demos above use the same underlying event — the crab feast, 42 of 44 replied, Bayside deposit paid, lake house committed at \$1,394 — so the comparison is of containers, not of content.

What they are actually using

	NIKE — "EXPERIENCES"	BLINK.GLOBAL — "EXPLORE"
Control above	Underline tabs — All · Customisation · Running · Training , scrollable, active = ink text + 2px underline	None. A section header only
Columns	2, equal	2, equal — twice: a featured pair, then a browse grid
Image ratio	~4:3, rounded	~4:3, rounded
Text placement	Below the image, always	Both. Below for entities, overlaid for categories
Text block	Bold title (wraps to 3 lines) · muted date · muted address	Featured: muted category + NAME UPPERCASE below · Browse: single uppercase word on the image
Overlay badge	None	Featured only — a gold ★ 8.4 top-left
Row heights	Uneven — content-driven	Even — fixed tile height

Read off the captured screens. Proportions are relative — Mobbin renders at roughly a third of native width, so no pixel value is asserted.

Correction to my earlier read

I said "Blink overlays" — that was half right

Blink runs **two different grids on one screen** and switches on what the tile represents: an *entity* (THE GRAND HOTEL, GREEN MEADOWS PARK) gets its text below the image; a *category* (CONFERENCE, MUSIC, EXPO) gets a single uppercase word on the image.

That is the better rule, and it is the one to carry: **overlay only survives when the label is one word you did not have to invent**. Our surface names are sentences, so the overlay branch is closed to us on every tile.

⚠ The constraint that decides this

Checked in the repo just now

Both grids are **photography-driven**. hostv2 ships exactly **two images** — `crab-hero.png` and `catfish-hero.png` — and both are event-type heroes. There is **no per-section artwork**.

so A literal port of either grid renders eight empty rectangles. My mock in the previous section used gradient placeholders and quietly implied images we do not have — the mocks below assume the real asset situation instead.

ALSO Commissioning eight section photos would collide with the artwork rule already in force: event glyphs are real artwork, identical at every size. Section photography is a different asset class and a standing cost.

Four mocks that work with what we have

All four use the real event — 42 of 44 replied, Bayside deposit paid, lake house committed at \$1,394, 9 days out. None needs a photograph.

A · NIKE'S SHAPE, NO PHOTOGRAPH

All Needs you Waiting Settled



Food & shopping



Where everyone
stays

B · THE STATE TILE — NUMBER LEADS

Fri

Crab count due

Bayside cutoff · 2
days

42 / 44

Replied

2 still out

Count due Friday ·
Bayside

Lake house ·
committed



Getting around

4 rides · 2
unassigned



Guests & RSVPs

42 of 44 replied

The image band becomes a quiet tinted cap holding a glyph. Keeps Nike's rhythm, needs no assets. Settled tiles drop to 72% opacity — state without colour.

\$1,394

Lake house

3 nights · committed

4

Rides

2 unassigned

Densest of the four and the most scannable — but it inverts the hierarchy: the number leads and the label follows. Only honest where a number is genuinely the answer.

C · DATA INSTEAD OF PHOTOGRAPHY



Run of show


Doors is next · 2 of 4
done



72%


Money

\$2,840 of \$3,950
committed



Vendors

3 booked · 1 awaiting
reply



Where everyone stays

Committed · backup
held

The missing photograph becomes an asset. Every bar here is computed already — `dayPhases()` returns `{total, done, state}`, plus `planned/committed/spent` and `raiseCounts`.

D · BLINK'S FEATURED SPLIT, HONESTLY

9 DAYS OUT

Crab feast · Saturday

Food & shopping

Count due Friday

Guests & RSVPs

42 of 44 replied

Where everyone stays

Committed

Getting around

2 unassigned

One hero using the one artwork we actually own — the event's own glyph — then unphotographed tiles beneath. Blink's structure, our asset reality.

C is the one I would build, with A's tab row above it. It solves the ragged problem structurally (equal tiles), needs no new assets, and turns the constraint into the feature — a Sections directory where every door states what it is holding. It also finally spends the infographic finding from the gap board, where the note was that the data is already computed in six places and rendered in one.

Caveat on C: one small visual per tile is four visuals per screen, which presses on the one-chart-per-surface rule. The rule was written for charts that carry a number; these carry a proportion and the number sits in the line beneath. Worth a ruling before building, not after.

Who is handling it best

SITUATION	WINNER	WHAT THEY DO
Several axes at once	Binance	Stacked labelled rails — the only app in the set that solves multi-axis without a modal
More options than fit	Hyundai Card	Rail + trailing  expander carrying the count of what is hidden
Long single list	komoot	Grouped sheet with an <i>All ...</i> row per group
Active filters must stay visible	Jobber	Anchored icon-pill with a removable X on each applied filter
Browse, not filter	blink.global, Nike	2-col tile grid — equal sizes, so nothing is ragged by construction
2–4 options	Nike	Underline tabs, equal width, no scroll at all
<i>Best-in-set by situation, from the screens read.</i>		

A spec you could gate

1. **Never wrap a pill row.** A test asserting the filter container has `flex-wrap: nowrap` and one computed line-height of children makes this permanent.
2. **2–4 options → segmented control.** Equal widths, all visible, no scroll.
3. **5–8 on one axis → scrolling rail,** mask the right edge so it clips mid-pill. Never fade to nothing at both ends unless both directions really have content.
4. **More than one axis → stack labelled rails.** One axis per row, 11px muted label above. Binance's pattern.
5. **9+ → grouped bottom sheet.** Radio right, label left, an *All ...* row per group.
6. **Selected = filled** (`--steel`), unselected = outlined on `--band` . Never rely on colour alone; the fill is the signal.
7. **If the rail exceeds 8, ask what to delete** before choosing a container.

Rules 2–5 are a single decision function of (count, axes), which means the container choice can be derived rather than argued about per surface.

Source: Mobbin iOS, read 1 August 2026 — Careem, Binance, Hyundai Card, komoot, Jobber, Nike, Zip, Crypto.com, Yubo, Formula 1, foodpanda, Meta AI, Booking.com, Wise, Journal, Azar, Linktree, KakaoTalk. Plus blink.global mobile frames, driven live.

Not claimed: *no hands-on with any of these apps, no exact paddings or type sizes — Mobbin renders at roughly a third of native width. The demos above are reconstructions in our own tokens, built to be draggable rather than to match their pixels. The "nobody wraps" claim covers the screens read here, not every screen those apps ship.*

Rendered in Studio Matte, single-theme by decision.